



INNOVATE

DATA AND AI/ML EDITION

March 9th, 2023

MLT04

Train ML models at scale with Amazon SageMaker

Ayman Salama

Sr. Partner Solution Architect, AWS

Agenda

Large Scale Training Challenges

Amazon SageMaker End-to-End Solution

Types of Distributed Training

Distributed Training features

Demo

Why do we need distributed training?

Models grow faster than hardware, leading to bottlenecks

INCREASING COMPLEXITY

- Businesses need higher precision in their model predictions
- Results in larger and more complex models
- Requires frequent retraining of models

INCREASING COSTS

- Increasing compute power required for frequent training of larger models drives up cost to train
- Becomes a barrier for innovation and growth


Compute power ~ 2x every 3.4 months

Deep and steep


Computing power used in training AI systems
Days spent calculating at one petaflop per second*, log scale

By fundamentals

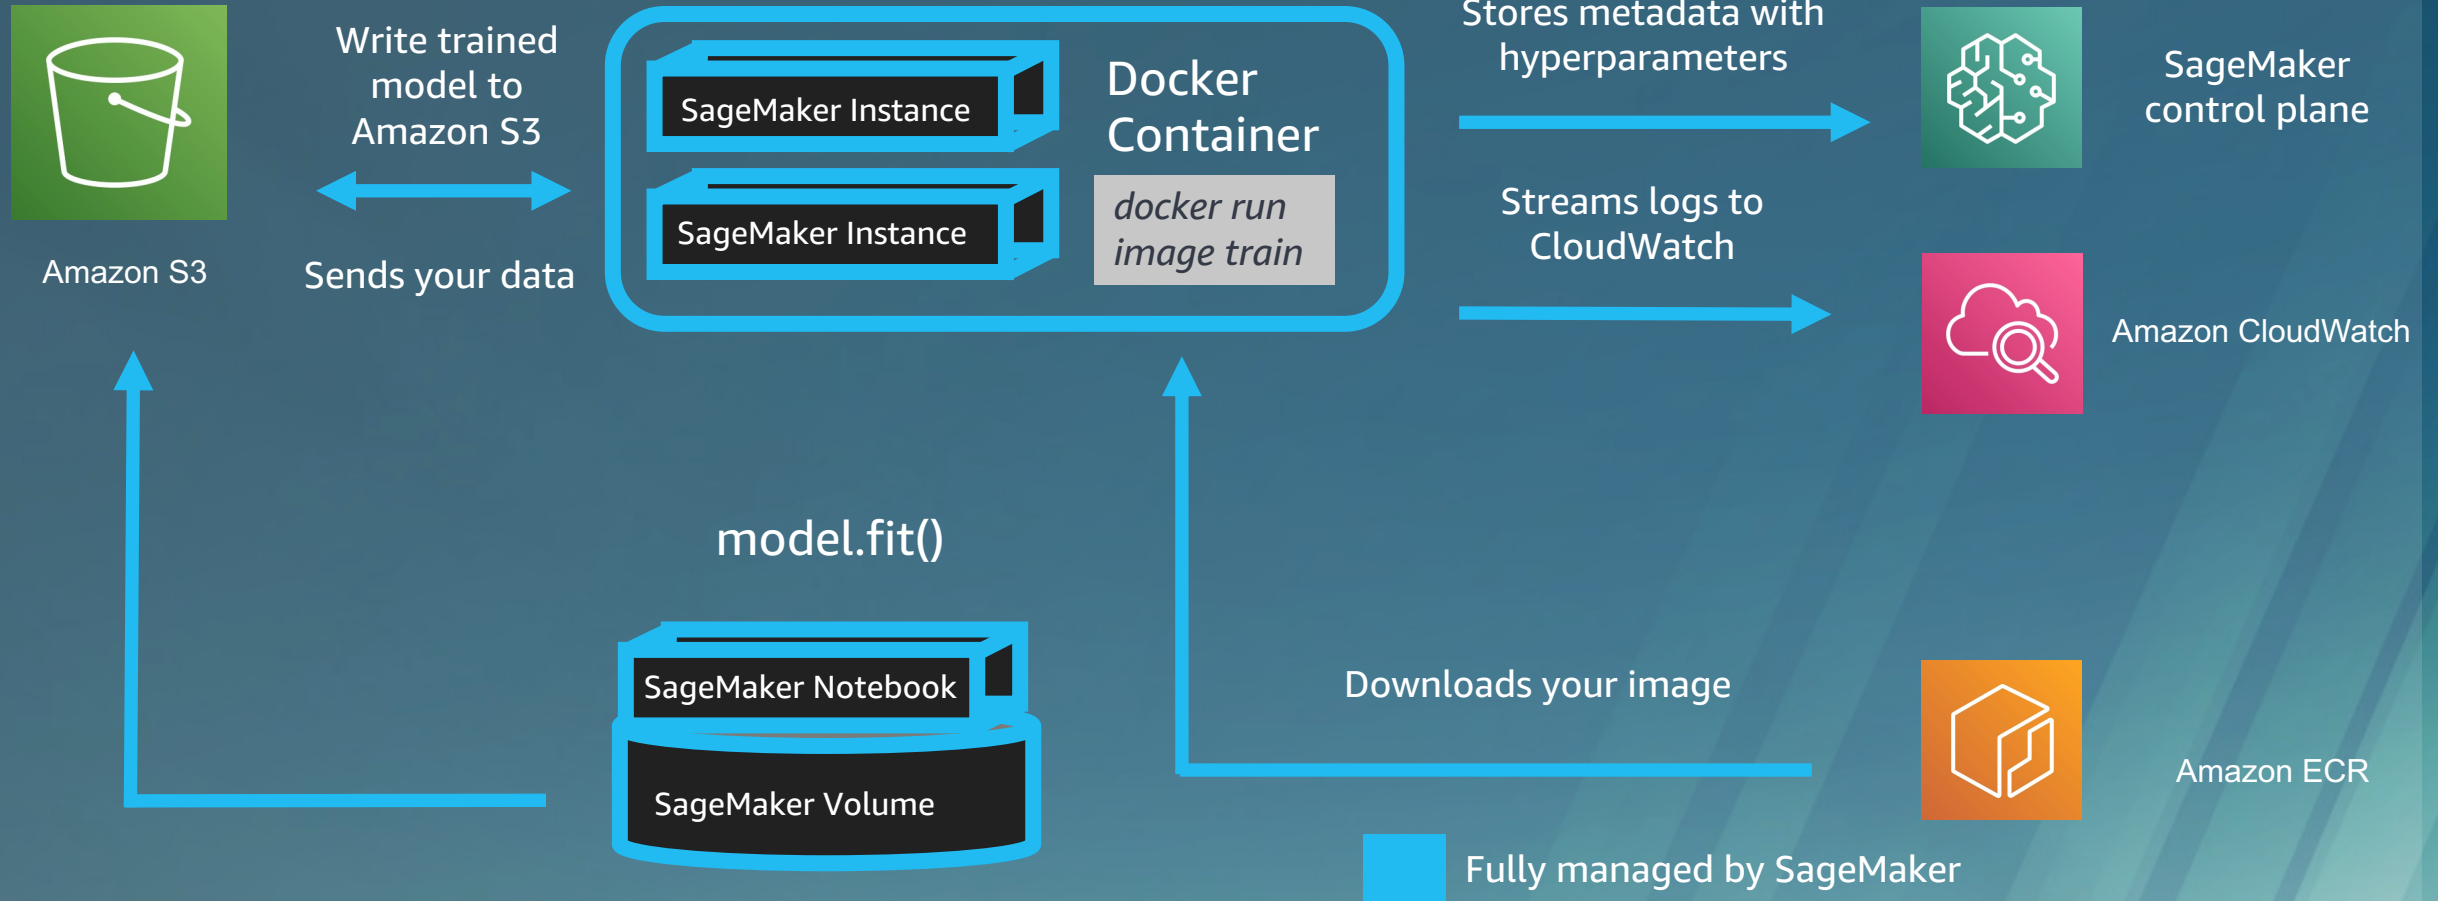
Language Speech Vision
Games Other



Model size increase ~ 10x/ year, Cost of Training increase ~ 100x by 2025



Amazon SageMaker ephemeral training clusters



ML.EC2 INSTANCES FOR DISTRIBUTED TRAINING

Broadest and most complete set of Distributed Training Infrastructure choices

G4dn

P3

P3dn

P4d

Ideal for

Accelerated training of small to medium sized models with less than 100M parameters

Training medium to large models with 100M to 300M parameters
Good for single node distributed training

Training large models with more than 300M parameters
Spot Training may provide better price-performance than P4d
Good for multi-node distributed training

Customer looking for best training performance on the cloud
Training large models with more than 300M parameters
Good for multi-node distributed training

Key features

16 GB/GPU

PCIe only

25 - 50 Gbps networking and 100 Gbps on bare-metal

16 GB/GPU

200 - 300 GB/s NVLink (4, 8 GPUs)

10-25 Gbps networking

32 GB/GPU

300 GB/s NVLink (8 GPUs)

100 Gbps networking

40 GB/GPU

600 GB/s NVLink (8 GPUs)

400 Gbps networking

GPU Config

1, 4, or 8 NVIDIA Tesla T4s

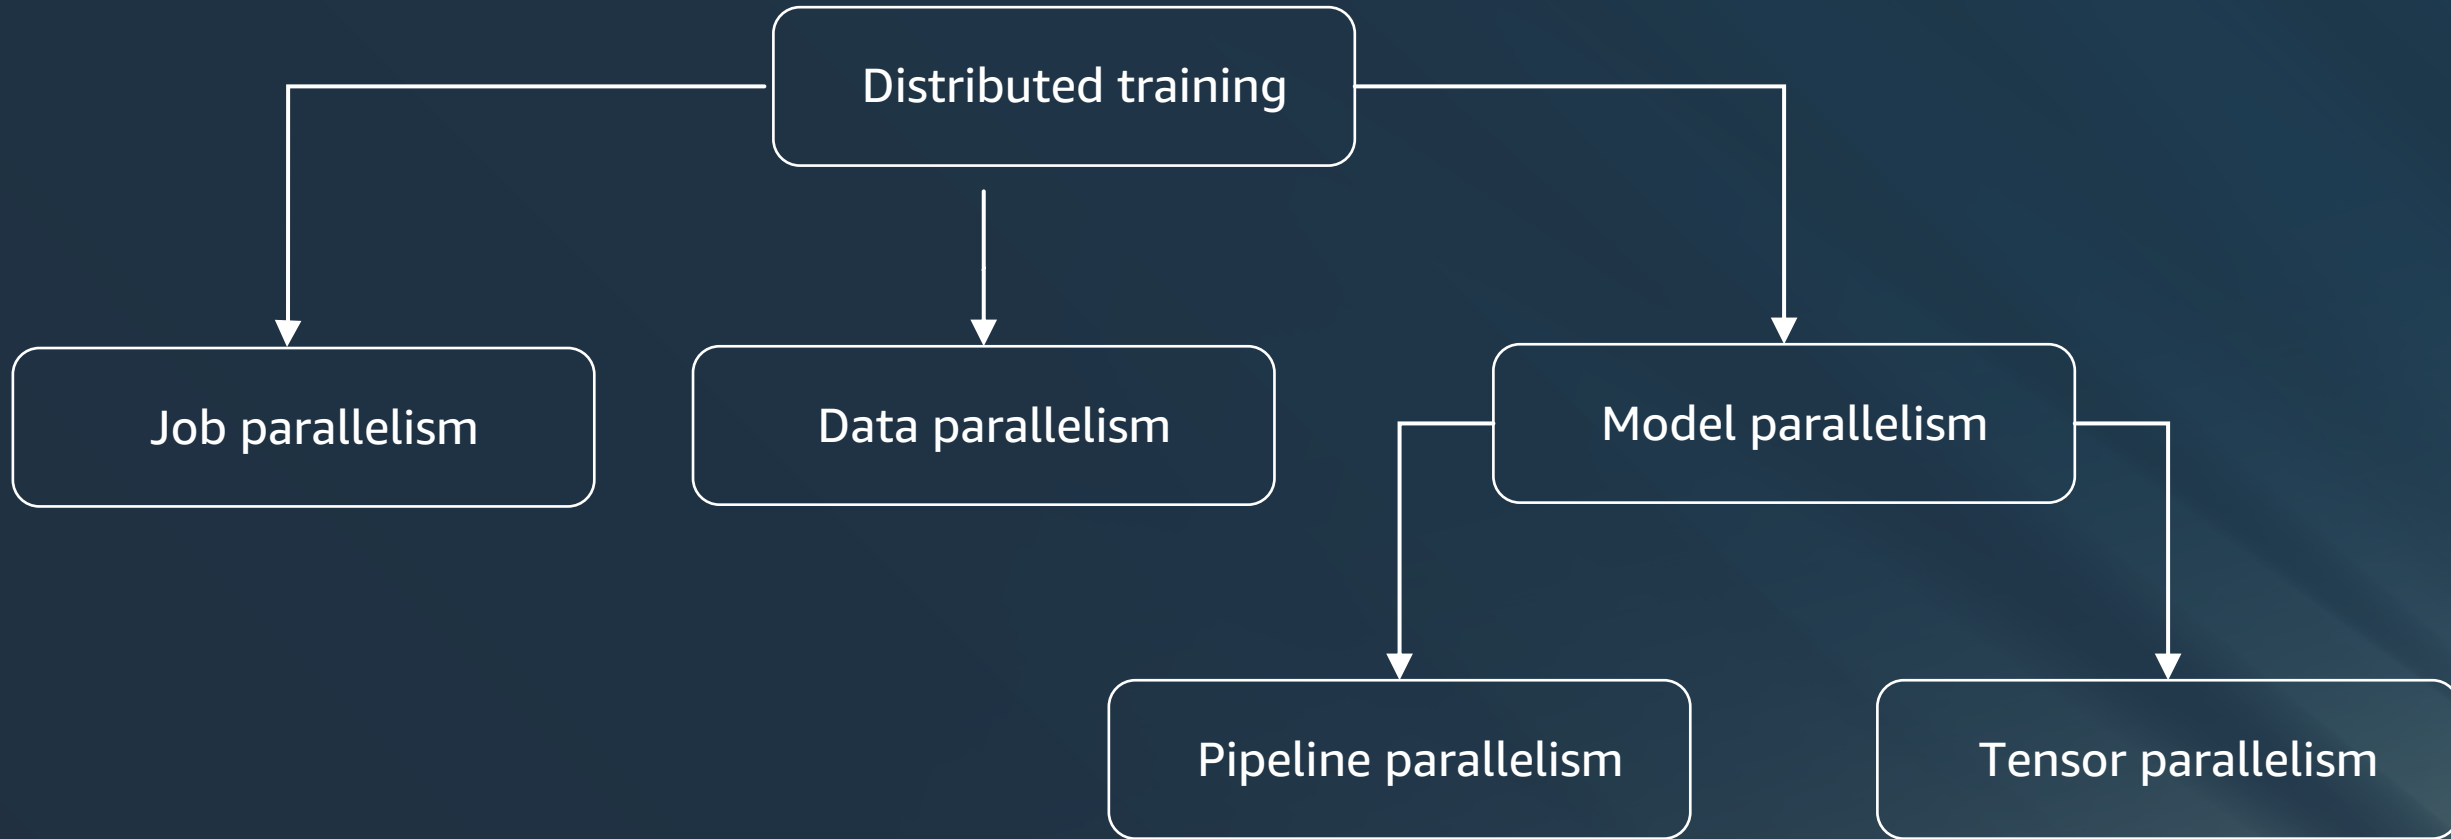
1, 4, or 8 NVIDIA Tesla V100s

8 NVIDIA Tesla V100s

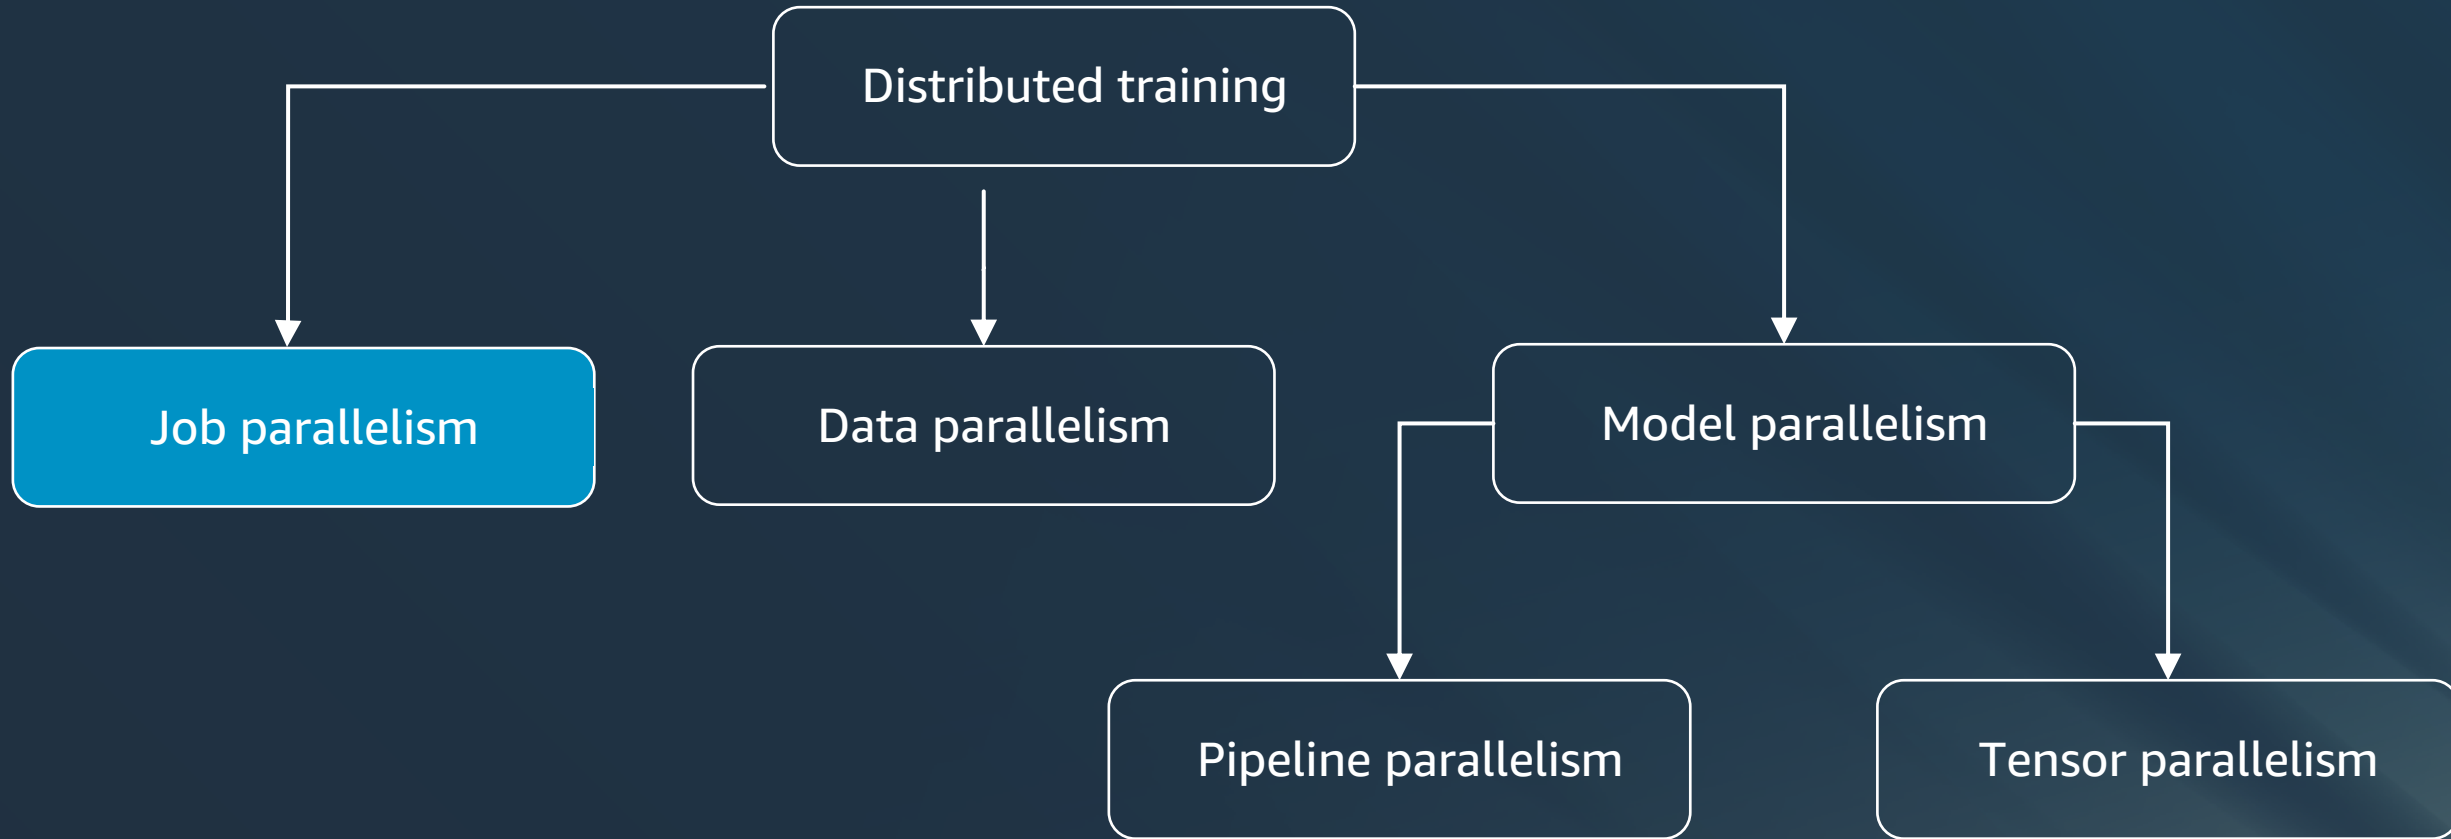
8 NVIDIA Tesla A100s (latest)



There are many kinds of distributed training



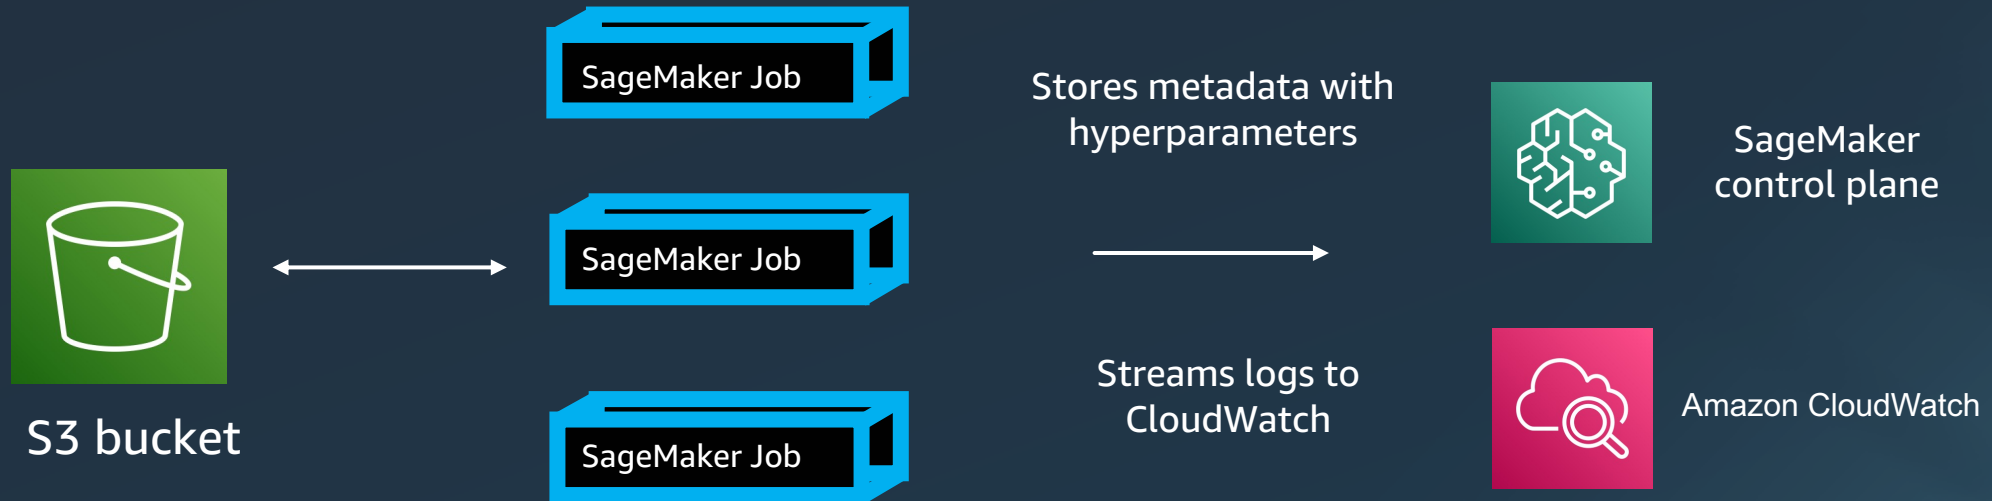
There are many kinds of distributed training



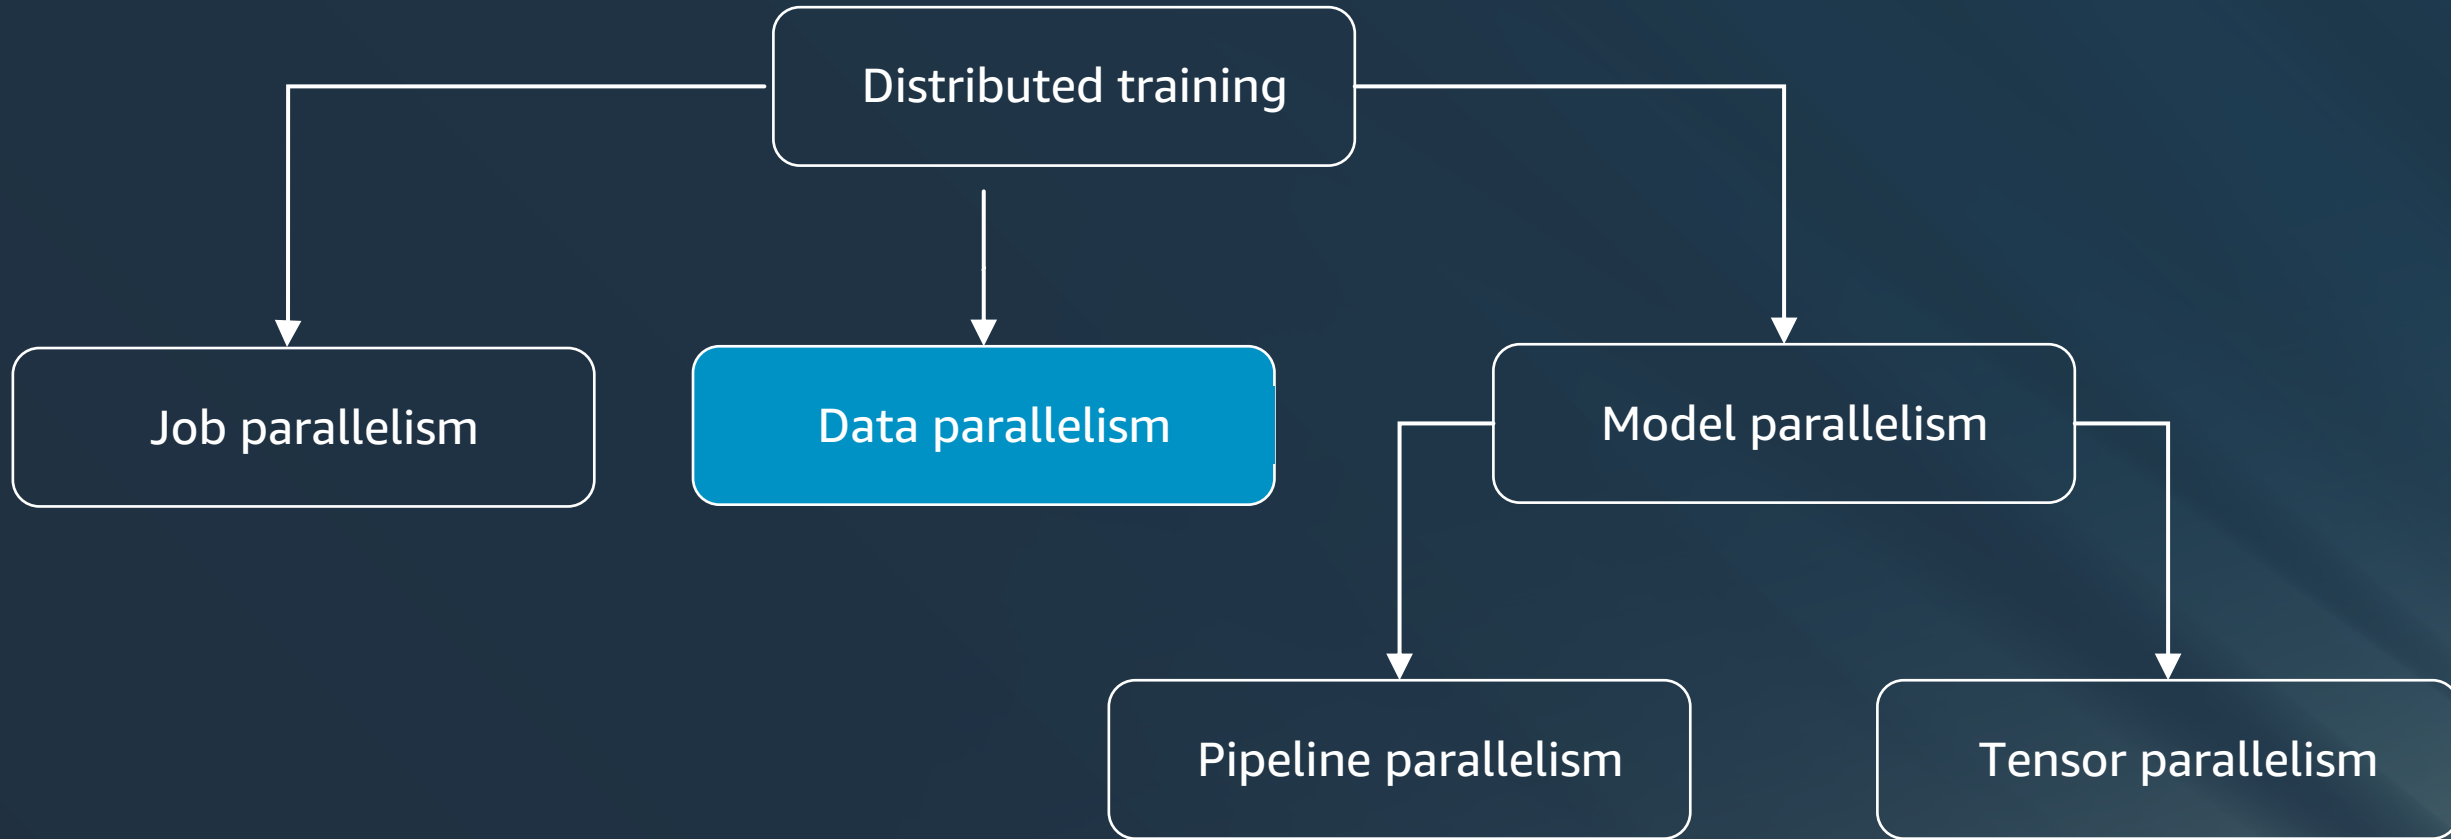
Train with parallel jobs at high frequency

1. Each job can train as many models as you need.
2. You can use *warm pools* to retrain as quickly as possible

```
for model in list_of_models:
    s3_input = get_data(model)
    s3_output = get_location(model)
    estimator = get_estimator(model, s3_output)
    estimator.fit(s3_input, wait=False)
```




There are many kinds of distributed training



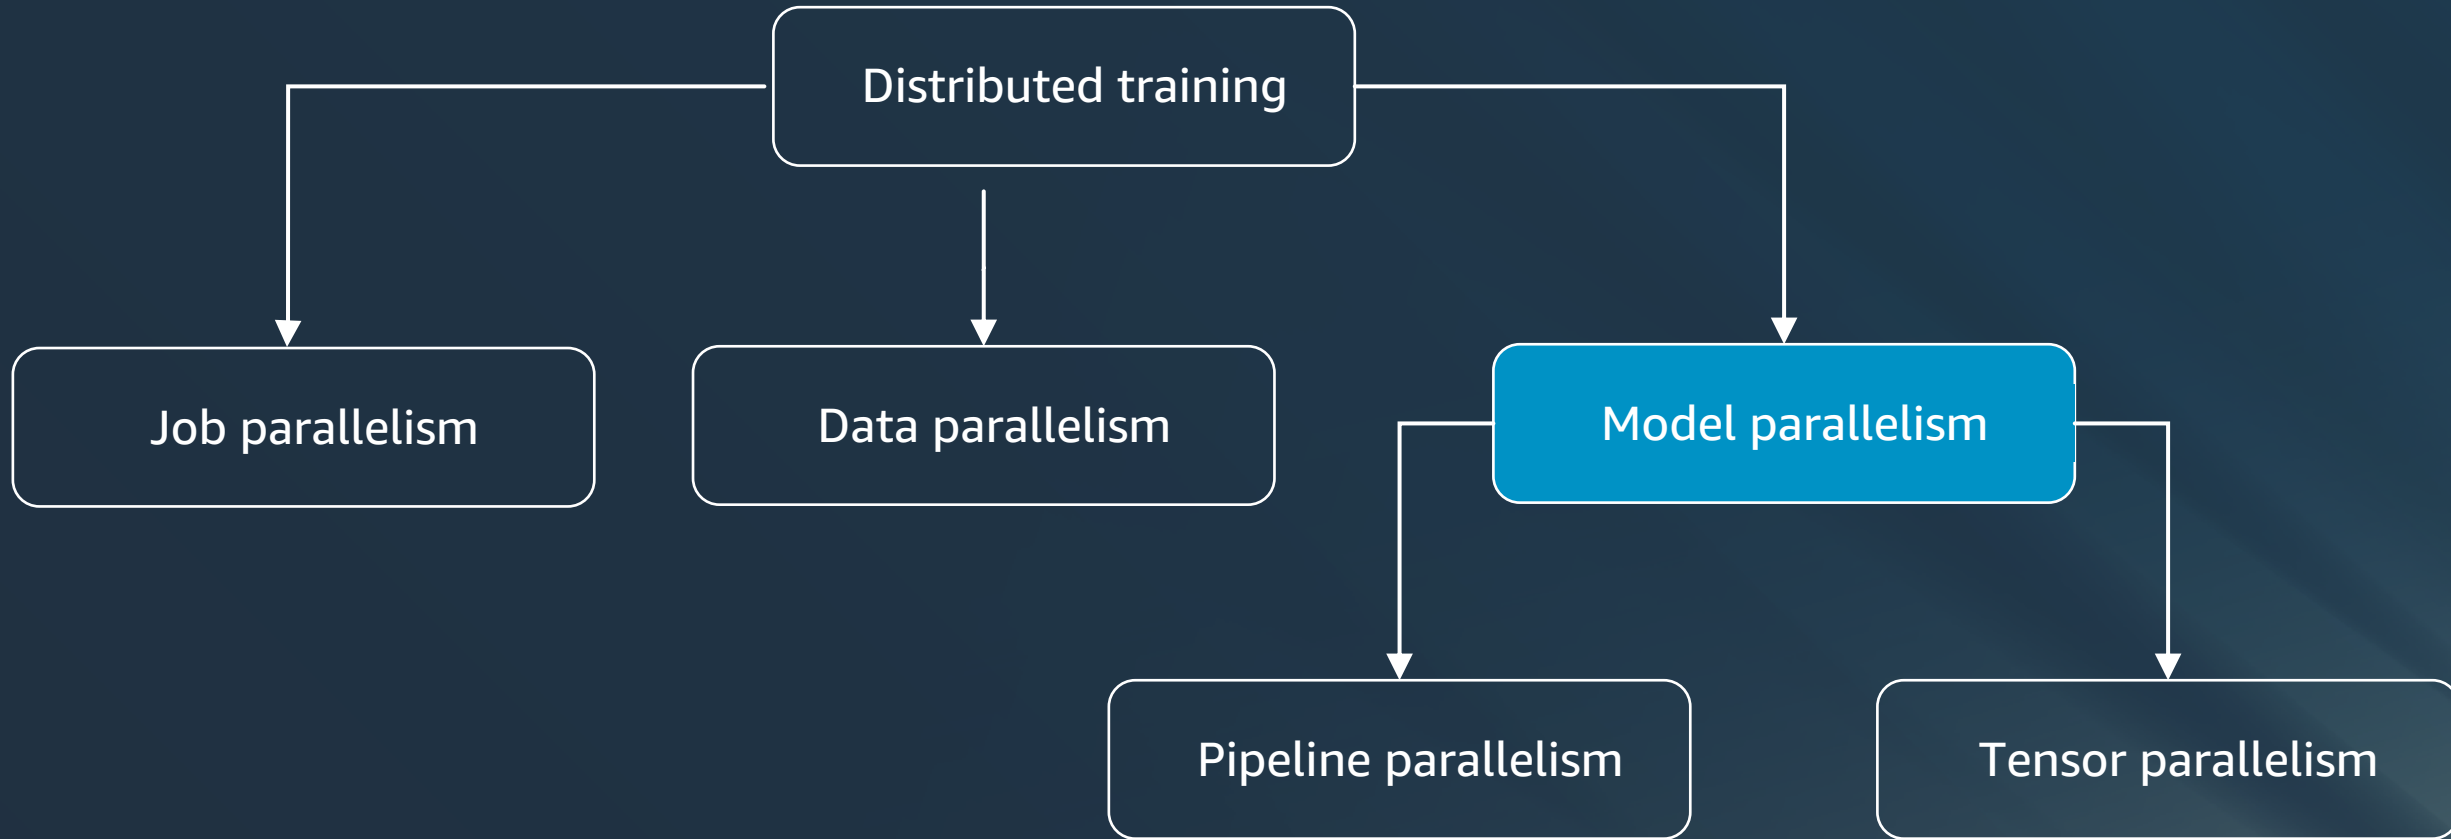
Data parallelism, think “massive data”

Amazon SageMaker Distributed Data Parallel

- **Optimized backend for distributed training of deep learning models in TensorFlow, PyTorch**
- **Accelerates training for network-bound workloads**
- **Built and optimized for AWS network topology and hardware**
- **20%–40% faster and cheaper than NCCL and MPI-based solutions. Best performance on AWS for large clusters.**




There are many kinds of distributed training



Model parallel, think “massive models”

Model parallelism on Amazon SageMaker (SMP)



**Automated
model partitioning**



**Interleaved
pipelined training**



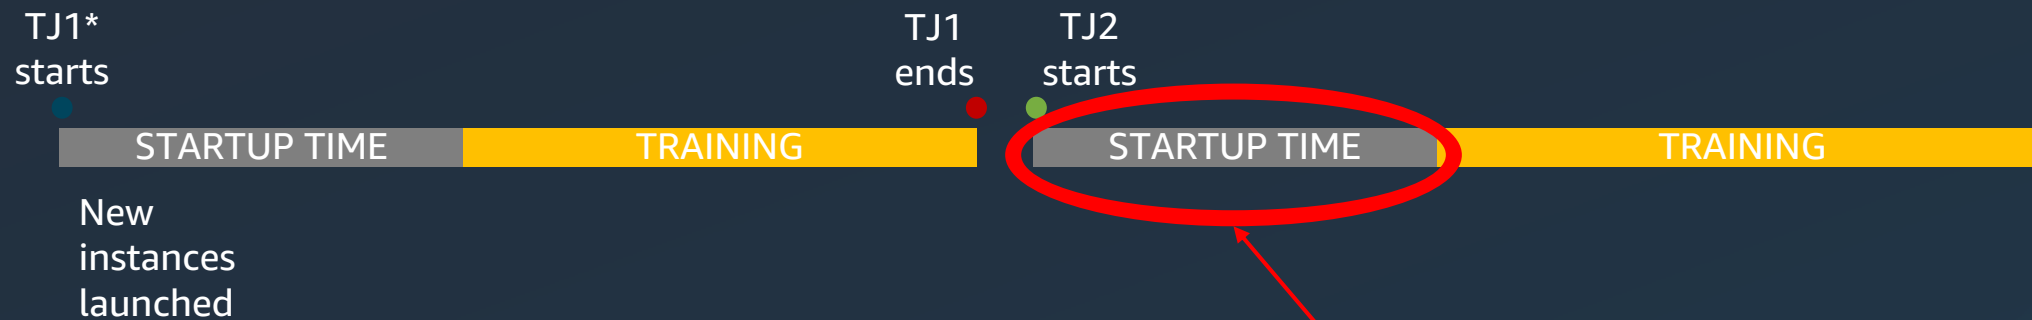
**Managed
SageMaker training**



**Clean
framework integration**

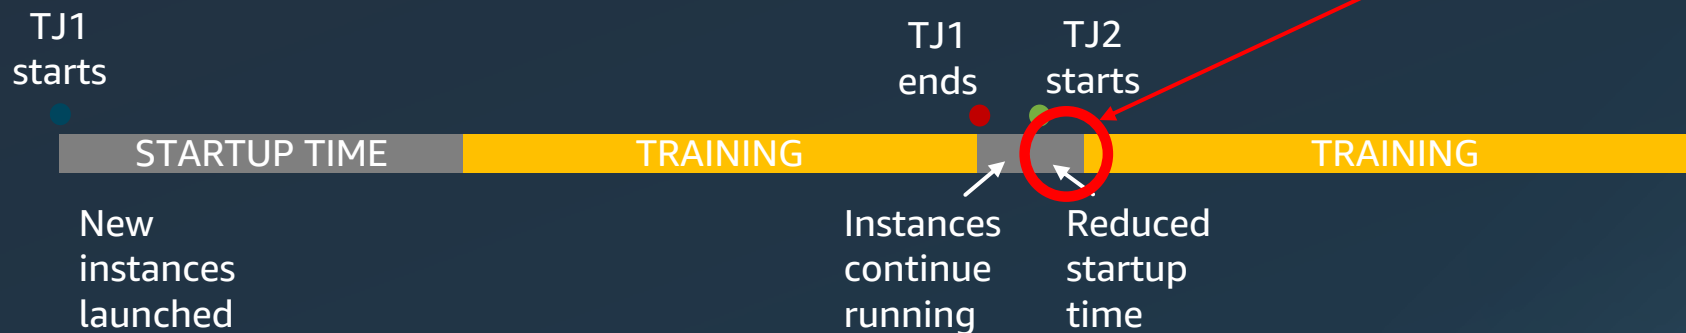
Warm Pools: Faster startup time

Before: Multi-minute wait between script updates



After: Multi-*second* wait between script updates

Keep_alive_period_in_seconds=600



Demo



+

Filter files by name

/ sagemaker-distributed-training-workshop / 2_model_parallel_sdp_Lab

Name

data_pipeline.py

data_prep_512.py

learning_rates.py

memory_tracker.py

requirements.txt

SageMaker-train-gpt-simple.ipynb

sharded_data_parallel_checkpoi...

train_gpt_simple.py

Home

Launcher

Terminal arn:aws:sagemaker:u

SageMaker-train-gpt-simple.ip

Home

Customize layout

Quick actions

Open Launcher

Create notebooks and other resources

Import & prepare data visually

Open the Getting Started notebook

Read documentation

View guided tutorials

Prebuilt and automated solutions

Deploy built-in algorithms, pre-built solutions, example notebooks, and build models from visual interface.

Quick start solutions

Pretrained models, notebooks, and prebuilt solutions

AutoML

Automatically build, train, and tune the best ML models

Workflows and tasks

Kick off a new step in the machine learning workflow.

Prepare data

Connect to data sources

Transform, analyze, and export data

Build, train, tune model

View all experiments

Create AutoML experiment

Deploy model

Get endpoint recommendation

Manage endpoints and optimize performance

Thank you!

Ayman Salama

aymanq@amazon.com

